

Active Accuracy Estimation on Large Datasets

Namit Katariya, Arun Iyer, Sunita Sarawagi

IIT Bombay

December 6, 2012

International Conference on Data Mining, 2012

Outline

- 1 Motivation
- 2 Problem Statement
- 3 Related Work
- 4 Proposed Solution
- 5 Results
- 6 Summary

Motivation

- Many applications rely on output of imperfect classifiers deployed on large datasets

Examples: Web page classification, classifying columns of a table to their semantic types

- **Common characteristics**

- Large and diverse dataset D
- Labeled data L unrepresentative of the entire dataset

- Measured accuracy on labeled set \neq True accuracy on data

- Need a method that can converge to the true accuracy

- ① An algorithm that returns a good estimate ($\hat{\mu}$) of true accuracy (μ) of the classifier
- ② *Scalable algorithm* : should work on large datasets where sequential scan not possible & data accessible only via an index

Motivation

- Many applications rely on output of imperfect classifiers deployed on large datasets

Examples: Web page classification, classifying columns of a table to their semantic types

- Common characteristics**

- Large and diverse dataset D
- Labeled data L unrepresentative of the entire dataset

- Measured accuracy on labeled set \neq True accuracy on data

- Need a method that can converge to the true accuracy

- 1 An algorithm that returns a good estimate ($\hat{\mu}$) of true accuracy (μ) of the classifier
- 2 *Scalable algorithm* : should work on large datasets where sequential scan not possible & data accessible only via an index

Motivation

- Many applications rely on output of imperfect classifiers deployed on large datasets

Examples: Web page classification, classifying columns of a table to their semantic types

- **Common characteristics**

- Large and diverse dataset D
 - Labeled data L unrepresentative of the entire dataset
- Measured accuracy on labeled set \neq True accuracy on data
- Need a method that can converge to the true accuracy
 - ① An algorithm that returns a good estimate ($\hat{\mu}$) of true accuracy (μ) of the classifier
 - ② *Scalable algorithm* : should work on large datasets where sequential scan not possible & data accessible only via an index

Motivation

- Many applications rely on output of imperfect classifiers deployed on large datasets

Examples: Web page classification, classifying columns of a table to their semantic types

- Common characteristics**

- Large and diverse dataset D
- Labeled data L unrepresentative of the entire dataset

- Measured accuracy on labeled set \neq True accuracy on data

- Need a method that can converge to the true accuracy

- ① An algorithm that returns a good estimate ($\hat{\mu}$) of true accuracy (μ) of the classifier
- ② *Scalable algorithm* : should work on large datasets where sequential scan not possible & data accessible only via an index

Motivation

- Many applications rely on output of imperfect classifiers deployed on large datasets

Examples: Web page classification, classifying columns of a table to their semantic types

- **Common characteristics**

- Large and diverse dataset D
 - Labeled data L unrepresentative of the entire dataset
- Measured accuracy on labeled set \neq True accuracy on data
- Need a method that can converge to the true accuracy
 - ① An algorithm that returns a good estimate ($\hat{\mu}$) of true accuracy (μ) of the classifier
 - ② *Scalable algorithm* : should work on large datasets where sequential scan not possible & data accessible only via an index

Outline

- 1 Motivation
- 2 Problem Statement**
- 3 Related Work
- 4 Proposed Solution
- 5 Results
- 6 Summary

Problem Statement

- ➊ **Accuracy estimation** : Estimate accuracy of a classifier on a large unlabeled dataset based on a **small, possibly unrepresentative, labeled set** and a human labeler
- ➋ **Scalable algorithm** : Perform accuracy estimation on unlabeled data so large that it makes **even a single sequential scan impractical** in an interactive setting

Outline

- 1 Motivation
- 2 Problem Statement
- 3 Related Work**
- 4 Proposed Solution
- 5 Results
- 6 Summary

Related Work

- Most existing work on *learning* rather than *evaluating* classifiers
- Existing works on selecting instances for evaluating classifiers:
 - (Sawade et al., 2010) present a new proposal distribution for sampling instances
 - (Bennett and Carvalho, 2010) and (Druck and McCallum, 2011) use stratified sampling

However, all assume classifier $C(\mathbf{x})$ to be probabilistic and use its $\Pr(y|\mathbf{x})$ scores for selection and/or stratification.

- Unlike (Bennett and Carvalho, 2010) and (Druck and McCallum, 2011), *instead of fixing the stratification, we learn a new one every time more data gets labeled*
- None of the existing methods consider cases where the dataset D is too large to even afford a single sequential scan

Related Work

- Most existing work on *learning* rather than *evaluating* classifiers
- Existing works on selecting instances for evaluating classifiers:
 - (Sawade et al., 2010) present a new proposal distribution for sampling instances
 - (Bennett and Carvalho, 2010) and (Druck and McCallum, 2011) use stratified sampling

However, all assume classifier $C(\mathbf{x})$ to be probabilistic and use its $\Pr(y|\mathbf{x})$ scores for selection and/or stratification.

- Unlike (Bennett and Carvalho, 2010) and (Druck and McCallum, 2011), *instead of fixing the stratification, we learn a new one every time more data gets labeled*
- None of the existing methods consider cases where the dataset D is too large to even afford a single sequential scan

Related Work

- Most existing work on *learning* rather than *evaluating* classifiers
- Existing works on selecting instances for evaluating classifiers:
 - (Sawade et al., 2010) present a new proposal distribution for sampling instances
 - (Bennett and Carvalho, 2010) and (Druck and McCallum, 2011) use stratified sampling

However, all assume classifier $C(\mathbf{x})$ to be probabilistic and use its $\Pr(y|\mathbf{x})$ scores for selection and/or stratification.

- Unlike (Bennett and Carvalho, 2010) and (Druck and McCallum, 2011), *instead of fixing the stratification, we learn a new one every time more data gets labeled*
- None of the existing methods consider cases where the dataset D is too large to even afford a single sequential scan

Related Work

- Most existing work on *learning* rather than *evaluating* classifiers
- Existing works on selecting instances for evaluating classifiers:
 - (Sawade et al., 2010) present a new proposal distribution for sampling instances
 - (Bennett and Carvalho, 2010) and (Druck and McCallum, 2011) use stratified sampling

However, all assume classifier $C(\mathbf{x})$ to be probabilistic and use its $\Pr(y|\mathbf{x})$ scores for selection and/or stratification.

- Unlike (Bennett and Carvalho, 2010) and (Druck and McCallum, 2011), *instead of fixing the stratification, we learn a new one every time more data gets labeled*
- None of the existing methods consider cases where the dataset D is too large to even afford a single sequential scan

Related Work

- Most existing work on *learning* rather than *evaluating* classifiers
- Existing works on selecting instances for evaluating classifiers:
 - (Sawade et al., 2010) present a new proposal distribution for sampling instances
 - (Bennett and Carvalho, 2010) and (Druck and McCallum, 2011) use stratified sampling

However, all assume classifier $C(\mathbf{x})$ to be probabilistic and use its $\Pr(y|\mathbf{x})$ scores for selection and/or stratification.

- Unlike (Bennett and Carvalho, 2010) and (Druck and McCallum, 2011), *instead of fixing the stratification, we learn a new one every time more data gets labeled*
- None of the existing methods consider cases where the dataset D is too large to even afford a single sequential scan

Outline

- 1 Motivation
- 2 Problem Statement
- 3 Related Work
- 4 Proposed Solution**
- 5 Results
- 6 Summary

Overall Idea

Algorithm 1 Loop for active accuracy estimation

- 1: $B = \# \text{buckets}$, $r = \# \text{bits}$, $\mathbf{f} = \text{feat. vector}$, $\mathbf{w}_{1\dots r} = \text{hyperplanes}$
 - 2: $\hat{\mu}_b, p_b = \text{accuracy \& weight estimates for bucket } b$
 - 3: **repeat**
 - 4: Learn stratification function $h(\mathbf{f}|\mathbf{w}_{1\dots r})$
 - 5: Stratify L via $h(\cdot)$ & compute $\{\hat{\mu}_b : 1 \leq b \leq B\}$
 - 6: Stratify D via $h(\cdot)$ & compute $\{p_b : 1 \leq b \leq B\}$
 - 7: Display accuracy estimates: $\hat{\mu}_S = \sum_b p_b \hat{\mu}_b$
 - 8: Get stratified sample set L' from D
 - 9: For each $\mathbf{x}_i \in L'$, get label y_i , and add (\mathbf{x}_i, y_i) to L
 - 10: **until** accuracy $\hat{\mu}_S$ not converged and labeler not bored.
 - 11: **Return** $\hat{\mu}_S$
-

Learning a stratification strategy

- Stratify input space so that instances in each stratum have similar accuracy values
 - *Supervised clustering methods* : Learn a distance function
Issue : Do not scale well
 - **Proposal** : Use **Hash codes** based on projections on hyperplanes learned over the feature space
- Learning hyperplanes (details in the paper)
 - *Smoothing the objective* : Convex upper bound
 - *Optimizing the smoothed objective* : EM-like algorithm
 - *Ensuring distinct hyperplanes* : Re-weight instances (boosting)
- $\hat{\mu}_b = \frac{1}{n_b} \sum_{i \in L_b} a_i$ prone to over-fitting. **Smooth based on labeled data in neighbouring buckets**
- Method agnostic to the type of classifier under consideration

Learning a stratification strategy

- Stratify input space so that instances in each stratum have similar accuracy values
 - *Supervised clustering methods* : Learn a distance function
Issue : Do not scale well
 - **Proposal** : Use **Hash codes** based on projections on hyperplanes learned over the feature space
- Learning hyperplanes (details in the paper)
 - *Smoothing the objective* : Convex upper bound
 - *Optimizing the smoothed objective* : EM-like algorithm
 - *Ensuring distinct hyperplanes* : Re-weight instances (boosting)
- $\hat{\mu}_b = \frac{1}{n_b} \sum_{i \in L_b} a_i$ prone to over-fitting. **Smooth based on labeled data in neighbouring buckets**
- Method agnostic to the type of classifier under consideration

Learning a stratification strategy

- Stratify input space so that instances in each stratum have similar accuracy values
 - *Supervised clustering methods* : Learn a distance function
Issue : Do not scale well
 - **Proposal** : Use **Hash codes** based on projections on hyperplanes learned over the feature space
- Learning hyperplanes (details in the paper)
 - *Smoothing the objective* : Convex upper bound
 - *Optimizing the smoothed objective* : EM-like algorithm
 - *Ensuring distinct hyperplanes* : Re-weight instances (boosting)
- $\hat{\mu}_b = \frac{1}{n_b} \sum_{i \in L_b} a_i$ prone to over-fitting. **Smooth based on labeled data in neighbouring buckets**
- Method agnostic to the type of classifier under consideration

Learning a stratification strategy

- Stratify input space so that instances in each stratum have similar accuracy values
 - *Supervised clustering methods* : Learn a distance function
Issue : Do not scale well
 - **Proposal** : Use **Hash codes** based on projections on hyperplanes learned over the feature space
- Learning hyperplanes (details in the paper)
 - *Smoothing the objective* : Convex upper bound
 - *Optimizing the smoothed objective* : EM-like algorithm
 - *Ensuring distinct hyperplanes* : Re-weight instances (boosting)
- $\hat{\mu}_b = \frac{1}{n_b} \sum_{i \in L_b} a_i$ prone to over-fitting. Smooth based on labeled data in neighbouring buckets
- Method agnostic to the type of classifier under consideration

Learning a stratification strategy

- Stratify input space so that instances in each stratum have similar accuracy values
 - *Supervised clustering methods* : Learn a distance function
Issue : Do not scale well
 - **Proposal** : Use **Hash codes** based on projections on hyperplanes learned over the feature space
- Learning hyperplanes (details in the paper)
 - *Smoothing the objective* : Convex upper bound
 - *Optimizing the smoothed objective* : EM-like algorithm
 - *Ensuring distinct hyperplanes* : Re-weight instances (boosting)
- $\hat{\mu}_b = \frac{1}{n_b} \sum_{i \in L_b} a_i$ prone to over-fitting. **Smooth based on labeled data in neighbouring buckets**
- Method agnostic to the type of classifier under consideration

Learning a stratification strategy

- Stratify input space so that instances in each stratum have similar accuracy values
 - *Supervised clustering methods* : Learn a distance function
Issue : Do not scale well
 - **Proposal** : Use **Hash codes** based on projections on hyperplanes learned over the feature space
- Learning hyperplanes (details in the paper)
 - *Smoothing the objective* : Convex upper bound
 - *Optimizing the smoothed objective* : EM-like algorithm
 - *Ensuring distinct hyperplanes* : Re-weight instances (boosting)
- $\hat{\mu}_b = \frac{1}{n_b} \sum_{i \in L_b} a_i$ prone to over-fitting. **Smooth based on labeled data in neighbouring buckets**
- Method agnostic to the type of classifier under consideration

Scaling up – Instance selection on large amounts of data

- Unlabeled data accessed for
 - ① computing p_b = the weight corresponding to each bucket b
 - ② generating sample L' from D to label and add to L
- Solutions for both, assigning bucket weights and selecting instances, based on sampling from a proposal distribution
- In each case, proposal distribution found by setting up an appropriate convex optimization problem
- Optimal proposal distribution can be calculated using some standard assumptions on the index (details in the paper)

Scaling up – Instance selection on large amounts of data

- Unlabeled data accessed for
 - ① computing p_b = the weight corresponding to each bucket b
 - ② generating sample L' from D to label and add to L
- Solutions for both, assigning bucket weights and selecting instances, based on sampling from a proposal distribution
- In each case, proposal distribution found by setting up an appropriate convex optimization problem
- Optimal proposal distribution can be calculated using some standard assumptions on the index (details in the paper)

Scaling up – Instance selection on large amounts of data

- Unlabeled data accessed for
 - ① computing p_b = the weight corresponding to each bucket b
 - ② generating sample L' from D to label and add to L
- Solutions for both, assigning bucket weights and selecting instances, based on sampling from a proposal distribution
- In each case, proposal distribution found by setting up an appropriate convex optimization problem
- Optimal proposal distribution can be calculated using some standard assumptions on the index (details in the paper)

Scaling up – Instance selection on large amounts of data

- Unlabeled data accessed for
 - ① computing p_b = the weight corresponding to each bucket b
 - ② generating sample L' from D to label and add to L
- Solutions for both, assigning bucket weights and selecting instances, based on sampling from a proposal distribution
- In each case, proposal distribution found by setting up an appropriate convex optimization problem
- Optimal proposal distribution can be calculated using some standard assumptions on the index (details in the paper)

Scaling up – Instance selection on large amounts of data

- Unlabeled data accessed for
 - ① computing p_b = the weight corresponding to each bucket b
 - ② generating sample L' from D to label and add to L
- Solutions for both, assigning bucket weights and selecting instances, based on sampling from a proposal distribution
- In each case, proposal distribution found by setting up an appropriate convex optimization problem
- Optimal proposal distribution can be calculated using some standard assumptions on the index (details in the paper)

Outline

- 1 Motivation
- 2 Problem Statement
- 3 Related Work
- 4 Proposed Solution
- 5 Results**
- 6 Summary

Results

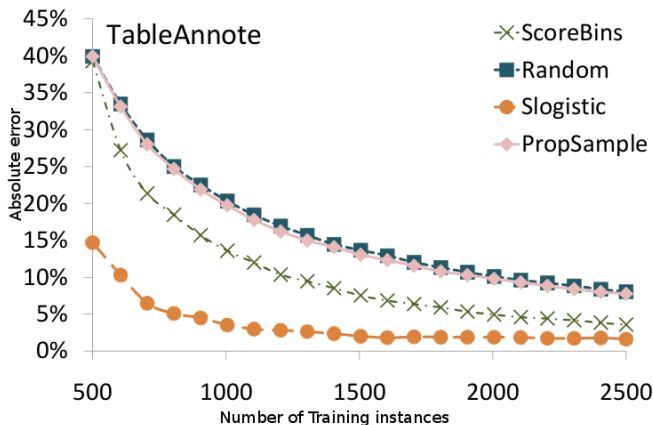
Summary of datasets used

- **TableAnnote** : Annotate columns of Web tables to type nodes of an ontology
- **Spam** : Classifying web-pages as spam or not
- **DNA** : Binary DNA classification task
- **HomeGround, HomePredicted** : Dataset of (entity, web-page) instances and decide if web-page was a homepage for the entity

Dataset	# Features	Size		Accuracy (%)	
		Seed(L)	Unlabeled(D)	Seed(L)	True(D)
TableAnnote	42	541	11,954,983	56.4	16.5
Spam	1000	5000	350,000	86.4	93.2
DNA	800	100,000	50,000,000	72.2	77.9
HomeGround	66	514	1060	50.4	32.8
HomePredicted	66	8658	13,951,053	83.2	93.9

Results

Comparison of estimation strategies on the TableAnnote dataset



Random	Random Sampling
PropSample	Sampling from proposal distribution
ScoreBins	Stratified sampling with scores

Results

Comparison of estimation strategies on remaining datasets

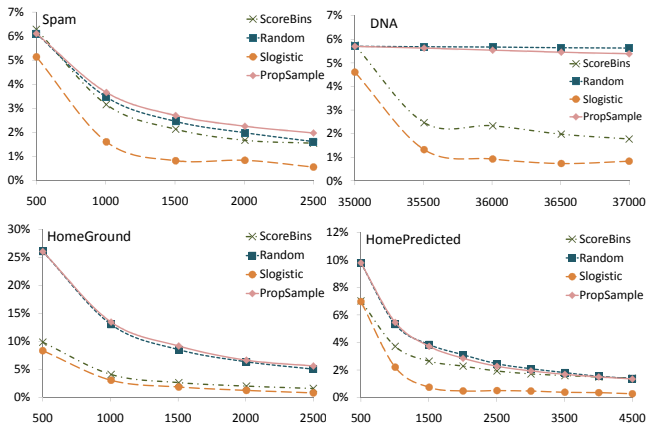
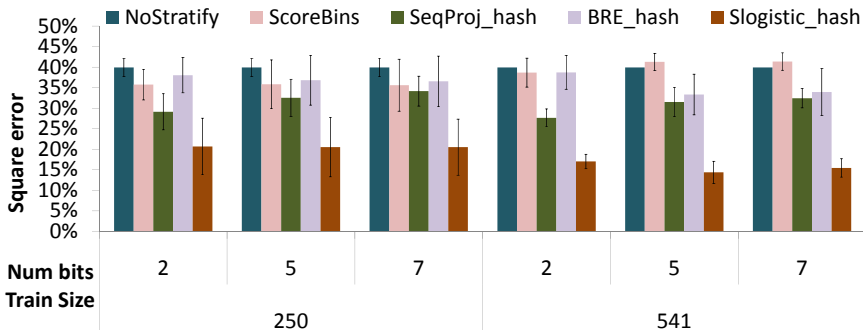


Figure: Absolute error (on the Y axis) of different estimation algorithms against increasing number of labeled instances (on the X axis)

Results

Comparison of different stratification methods on the TableAnnote dataset

TableAnnote



NoStratify	Simple averaging (no stratification)
ScoreBins	Stratify using classifier scores
SeqProj_hash	Learn hyperplanes using method in (Wang et al., 2010)
BRE_hash	Learn hyperplanes using method in (Kulis and Darrell, 2009)

Results

Comparison of different stratification methods on remaining datasets

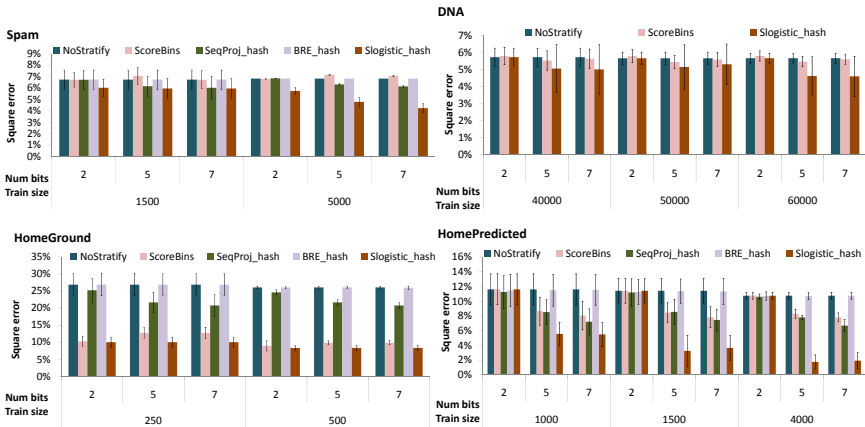


Figure: Error of different stratification methods against increasing training sizes and for different number of bits

Results

Comparison of methods of sampling from indexed data for estimating bucket weights

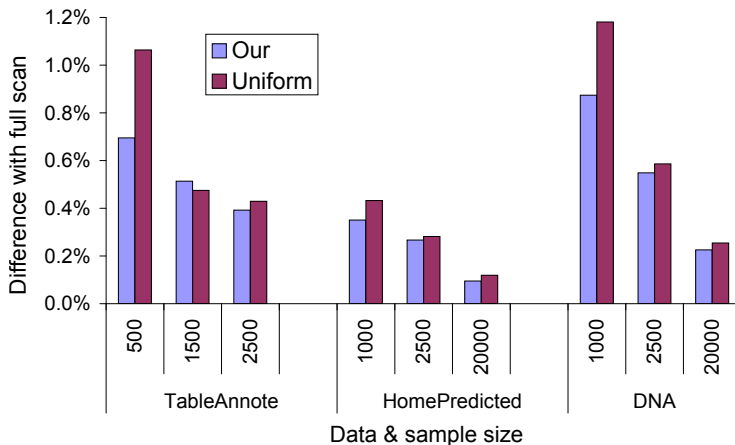


Figure: Comparing methods of sampling from indexed data for estimating bucket weights

Outline

- 1 Motivation
- 2 Problem Statement
- 3 Related Work
- 4 Proposed Solution
- 5 Results
- 6 Summary**

Summary

- 1 Addressed the challenge of *calibrating a classifier's accuracy on large unlabeled datasets* given small amounts of labeled data and a human labeler
- 2 Proposed a stratified sampling-based method for accuracy estimation that provides better estimates than simple averaging & better selection of instances for labeling than random sampling
- 3 Between 15% and 62% relative reduction in error achieved compared to existing approaches
- 4 Algorithm made *scalable* by proposing optimal sampling strategies for accessing indexed unlabeled data directly
- 5 Close to optimal performance while reading three orders of magnitude fewer instances on large datasets

Summary

- 1 Addressed the challenge of *calibrating a classifier's accuracy on large unlabeled datasets* given small amounts of labeled data and a human labeler
- 2 Proposed a stratified sampling-based method for accuracy estimation that provides better estimates than simple averaging & better selection of instances for labeling than random sampling
- 3 Between 15% and 62% relative reduction in error achieved compared to existing approaches
- 4 Algorithm made *scalable* by proposing optimal sampling strategies for accessing indexed unlabeled data directly
- 5 Close to optimal performance while reading three orders of magnitude fewer instances on large datasets

Summary

- 1 Addressed the challenge of *calibrating a classifier's accuracy on large unlabeled datasets* given small amounts of labeled data and a human labeler
- 2 Proposed a stratified sampling-based method for accuracy estimation that provides better estimates than simple averaging & better selection of instances for labeling than random sampling
- 3 Between 15% and 62% relative reduction in error achieved compared to existing approaches
- 4 Algorithm made *scalable* by proposing optimal sampling strategies for accessing indexed unlabeled data directly
- 5 Close to optimal performance while reading three orders of magnitude fewer instances on large datasets

Summary

- ① Addressed the challenge of *calibrating a classifier's accuracy on large unlabeled datasets* given small amounts of labeled data and a human labeler
- ② Proposed a stratified sampling-based method for accuracy estimation that provides better estimates than simple averaging & better selection of instances for labeling than random sampling
- ③ Between 15% and 62% relative reduction in error achieved compared to existing approaches
- ④ Algorithm made *scalable* by proposing optimal sampling strategies for accessing indexed unlabeled data directly
- ⑤ Close to optimal performance while reading three orders of magnitude fewer instances on large datasets

Summary

- ① Addressed the challenge of *calibrating a classifier's accuracy on large unlabeled datasets* given small amounts of labeled data and a human labeler
- ② Proposed a stratified sampling-based method for accuracy estimation that provides better estimates than simple averaging & better selection of instances for labeling than random sampling
- ③ Between 15% and 62% relative reduction in error achieved compared to existing approaches
- ④ Algorithm made *scalable* by proposing optimal sampling strategies for accessing indexed unlabeled data directly
- ⑤ Close to optimal performance while reading three orders of magnitude fewer instances on large datasets

Thank You

References I

- Bennett, P. N. and Carvalho, V. R. (2010). Online stratified sampling: evaluating classifiers at web-scale. In *CIKM*.
- Druck, G. and McCallum, A. (2011). Toward interactive training and evaluation. In *CIKM*.
- Kulis, B. and Darrell, T. (2009). Learning to hash with binary reconstructive embeddings. In *NIPS*.
- Sawade, C., Landwehr, N., Bickel, S., and Scheffer, T. (2010). Active risk estimation. In *ICML*.
- Wang, J., Kumar, S., and Chang, S. (2010). Sequential projection learning for hashing with compact codes. In *ICML*.

Assigning Bucket Weights

- Sample from a proposal distribution : $\hat{\mu}_{S_q} = \frac{1}{|S|} \sum_{\mathbf{x} \in S} \frac{1/N}{q(\mathbf{x})} \hat{\mu}_{h(\mathbf{x})}$
- **Result** : When $q(\mathbf{x})$ is restricted so that all instances within a partition u are sampled with the same probability q_u , the expected squared error between $\hat{\mu}_{S_q}$ and $\hat{\mu}_S$ is minimized when

$$q_u \propto \sqrt{\sum_b \hat{\mu}_b^2 p(b|u)}$$

- $p(b|u)$ = fraction of instances in D_u with $h(\mathbf{x}) = b$
- Initially, use labeled data to estimate $p(b|u)$
- As more instances are sampled from any D_u , refine estimates of $p(b|u)$

Instance Selection

- Perform importance sampling where $\text{imp}(\mathbf{x}) \propto \hat{\sigma}_{h(\mathbf{x})}$ without evaluating $h(\mathbf{x})$ over entire D
- Generate a larger sample S via proposal distribution $q(\mathbf{x})$ restricted to choose same $q(\mathbf{x}) \forall \mathbf{x}$ in data partition D_u
- Then from S generate the sample of k instances by weighting each instance as $f(\mathbf{x})/q(\mathbf{x})$. Good only if $q(\mathbf{x}) \sim f(\mathbf{x})$
- Best $q(\mathbf{x})$ found by solving for unlabeled bucket weights q_1, \dots, q_U so that expected L1 distance between $f(\mathbf{x})$ and $q(\mathbf{x})$ is minimized

$$\min_{q_1, \dots, q_U} \sum_u \sum_b p_u p(b|u) \left| \frac{\hat{\sigma}_b}{Z_f} - q_u \right| \text{ s.t. } \sum_u N p_u q_u = 1$$

- Z_f approximated as $\sum_b \hat{\sigma}_b \sum_u p_u p(b|u)$
- Get $p(b|u)$ as explained in the previous slide